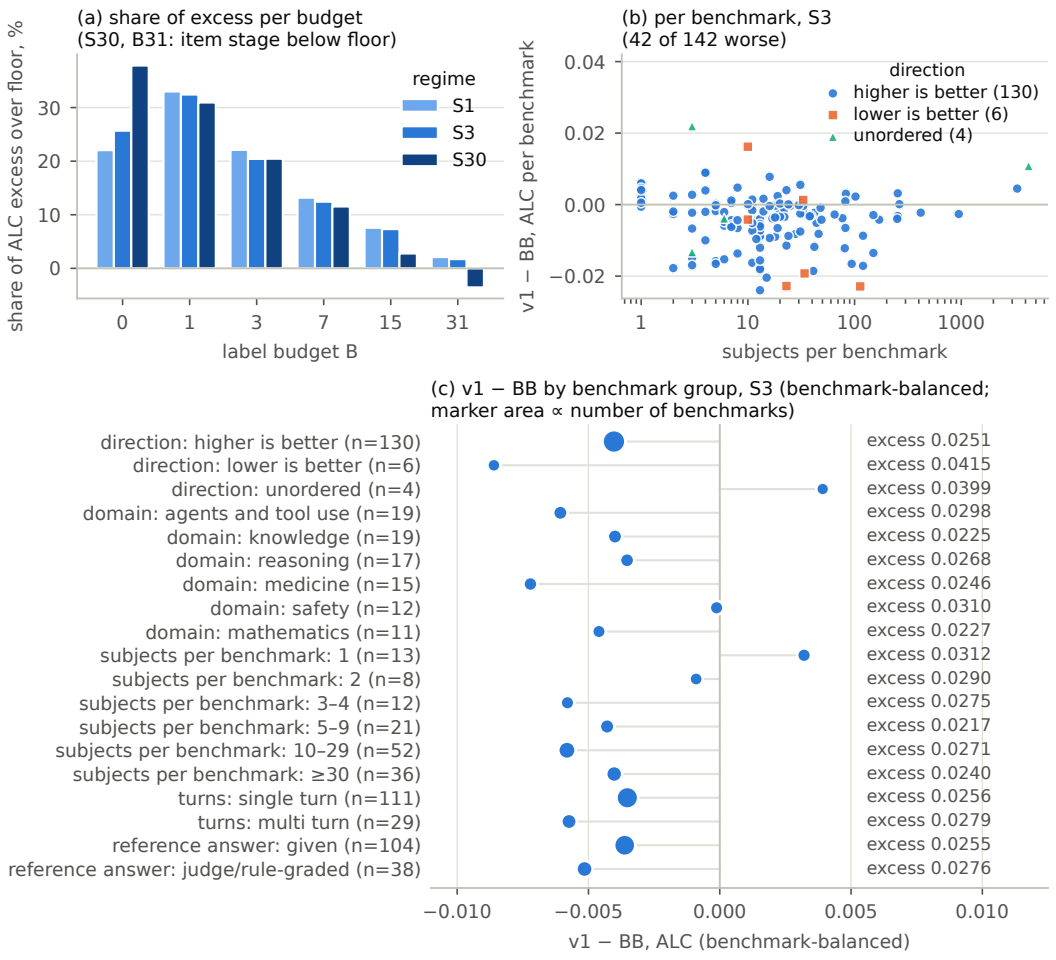
(a) share of excess per budget
(S30, B31: item stage below floor)
share of ALC excess over floor, %
regime
S1
S3
S30
0
10
20
30
0
1
3
7
15
31
label budget B
(b) per benchmark, S3
(42 of 142 worse)
v1 − BB, ALC per benchmark
direction
higher is better (130)
lower is better (6)
unordered (4)
0.04
0.02
0.00
−0.02
1
10
100
1000
subjects per benchmark
(c) v1 − BB by benchmark group, S3 (benchmark-balanced;
marker area ∝ number of benchmarks)
direction: higher is better (n=130)
direction: lower is better (n=6)
direction: unordered (n=4)
domain: agents and tool use (n=19)
domain: knowledge (n=19)
domain: reasoning (n=17)
domain: medicine (n=15)
domain: safety (n=12)
domain: mathematics (n=11)
subjects per benchmark: 1 (n=13)
subjects per benchmark: 2 (n=8)
subjects per benchmark: 3–4 (n=12)
subjects per benchmark: 5–9 (n=21)
subjects per benchmark: 10–29 (n=52)
subjects per benchmark: ≥30 (n=36)
turns: single turn (n=111)
turns: multi turn (n=29)
reference answer: given (n=104)
reference answer: judge/rule-graded (n=38)
excess 0.0251
excess 0.0415
excess 0.0399
excess 0.0298
excess 0.0225
excess 0.0268
excess 0.0246
excess 0.0310
excess 0.0227
excess 0.0312
excess 0.0290
excess 0.0275
excess 0.0217
excess 0.0271
excess 0.0240
excess 0.0256
excess 0.0279
excess 0.0255
excess 0.0276
−0.010
−0.005
0.000
0.005
0.010
v1 − BB, ALC (benchmark-balanced)