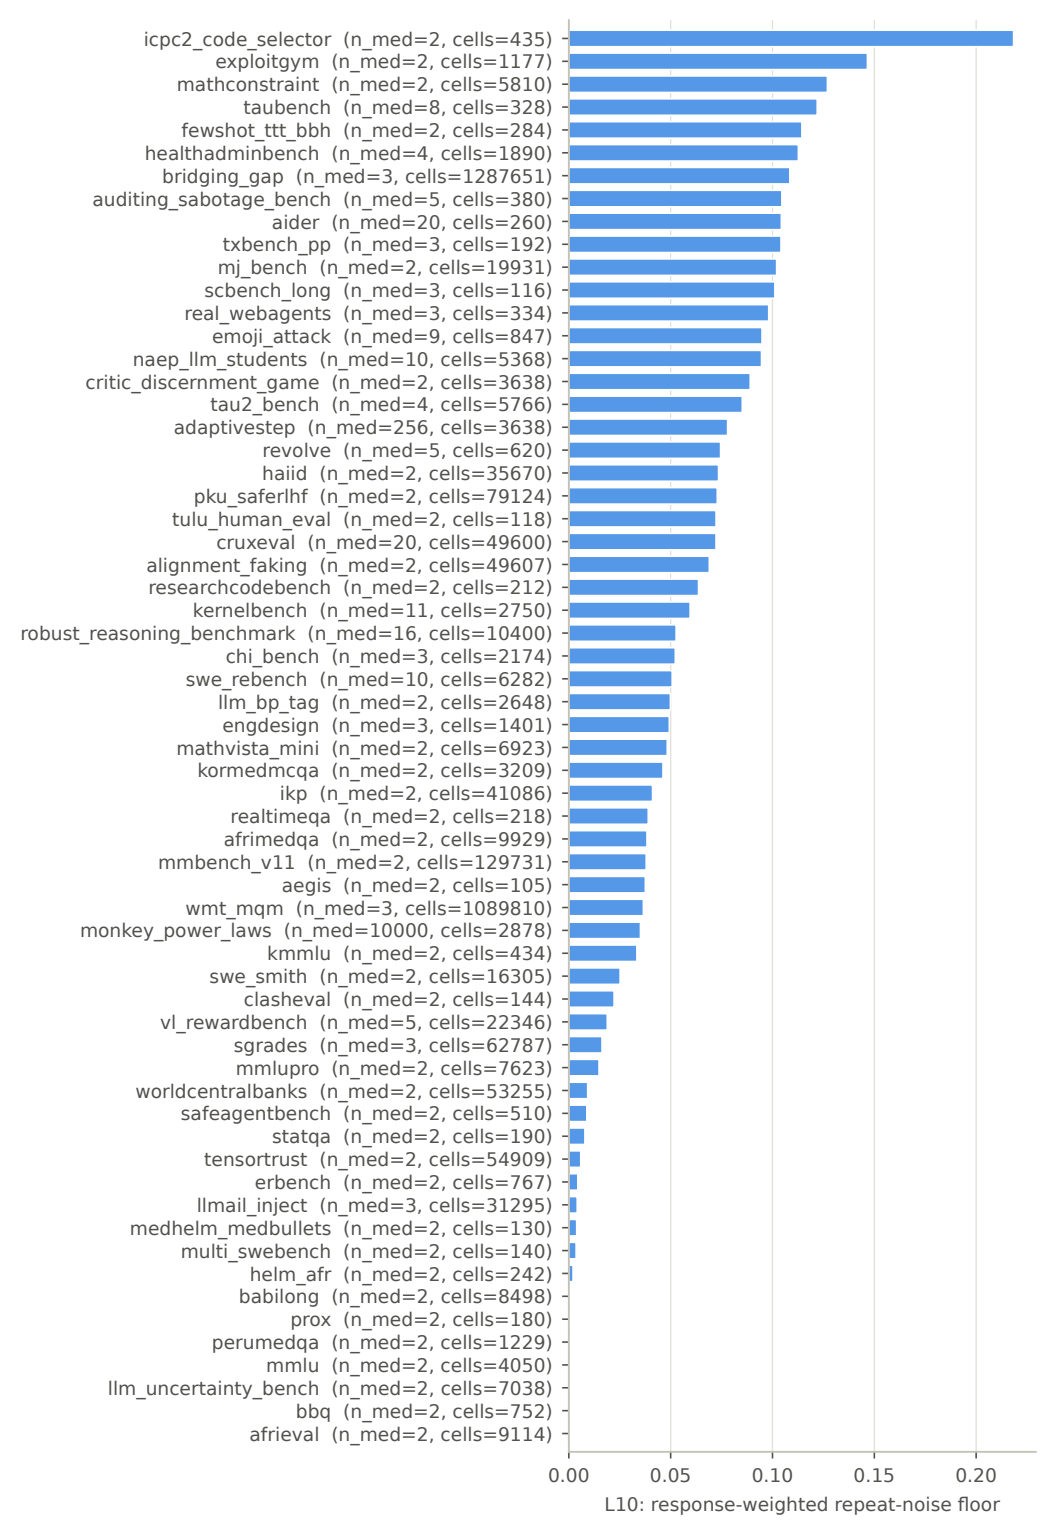
icpc2_code_selector (n_med=2, cells=435)
exploitgym (n_med=2, cells=1177)
mathconstraint (n_med=2, cells=5810)
taubench (n_med=8, cells=328)
fewshot_ttt_bbh (n_med=2, cells=284)
healthadminbench (n_med=4, cells=1890)
bridging_gap (n_med=3, cells=1287651)
auditing_sabotage_bench (n_med=5, cells=380)
aider (n_med=20, cells=260)
txbench_pp (n_med=3, cells=192)
mj_bench (n_med=2, cells=19931)
scbench_long (n_med=3, cells=116)
real_webagents (n_med=3, cells=334)
emoji_attack (n_med=9, cells=847)
naep_llm_students (n_med=10, cells=5368)
critic_discernment_game (n_med=2, cells=3638)
tau2_bench (n_med=4, cells=5766)
adaptivestep (n_med=256, cells=3638)
revolve (n_med=5, cells=620)
haiid (n_med=2, cells=35670)
pku_saferlhf (n_med=2, cells=79124)
tulu_human_eval (n_med=2, cells=118)
cruxeval (n_med=20, cells=49600)
alignment_faking (n_med=2, cells=49607)
researchcodebench (n_med=2, cells=212)
kernelbench (n_med=11, cells=2750)
robust_reasoning_benchmark (n_med=16, cells=10400)
chi_bench (n_med=3, cells=2174)
swe_rebench (n_med=10, cells=6282)
llm_bp_tag (n_med=2, cells=2648)
engdesign (n_med=3, cells=1401)
mathvista_mini (n_med=2, cells=6923)
kormedmcqa (n_med=2, cells=3209)
ikp (n_med=2, cells=41086)
realtimeqa (n_med=2, cells=218)
afrimedqa (n_med=2, cells=9929)
mmbench_v11 (n_med=2, cells=129731)
aegis (n_med=2, cells=105)
wmt_mqm (n_med=3, cells=1089810)
monkey_power_laws (n_med=10000, cells=2878)
kmmlu (n_med=2, cells=434)
swe_smith (n_med=2, cells=16305)
clasheval (n_med=2, cells=144)
vl_rewardbench (n_med=5, cells=22346)
sgrades (n_med=3, cells=62787)
mmlupro (n_med=2, cells=7623)
worldcentralbanks (n_med=2, cells=53255)
safeagentbench (n_med=2, cells=510)
statqa (n_med=2, cells=190)
tensortrust (n_med=2, cells=54909)
erbench (n_med=2, cells=767)
llmail_inject (n_med=3, cells=31295)
medhelm_medbullets (n_med=2, cells=130)
multi_swebench (n_med=2, cells=140)
helm_afr (n_med=2, cells=242)
babilong (n_med=2, cells=8498)
prox (n_med=2, cells=180)
perumedqa (n_med=2, cells=1229)
mmlu (n_med=2, cells=4050)
llm_uncertainty_bench (n_med=2, cells=7038)
bbq (n_med=2, cells=752)
afrieval (n_med=2, cells=9114)
0.00
0.05
0.10
0.15
0.20
L10: response-weighted repeat-noise floor